

Manipulating Data in R

Introduction to R for Public Health Researchers

Reshaping Data

In this module, we will show you how to:

1. Reshaping data from wide (fat) to long (tall)
2. Reshaping data from long (tall) to wide (fat)
3. Merging Data
4. Perform operations by a grouping variable

Setup

We will show you how to do each operation in base R then show you how to use the `dplyr` or `tidyr` package to do the same operation (if applicable).

See the “Data Wrangling Cheat Sheet using `dplyr` and `tidyr`”:

- ▶ <https://www.rstudio.com/wp-content/uploads/2015/02/data-wrangling-cheatsheet.pdf>

What is wide/long data?

See http://www.cookbook-r.com/Manipulating_data/Converting_data_between_wide_and_long_format/

- ▶ Wide - multiple columns per observation
 - ▶ e.g. visit1, visit2, visit3

	id	visit1	visit2	visit3
1	1	10	4	3
2	2	5	6	NA

- ▶ Long - multiple rows per observation

	id	visit	value
1	1	1	10
2	1	2	4
3	1	3	3
4	2	1	5
5	2	2	6

What is wide/long data?

More accurately, data is wide or long **with respect** to certain variables.

Data used: Charm City Circulator

`http://www.aejaffe.com/winterR_2017/data/Charm_City_Circulator_Ridership.csv`

```
circ = read_csv(  
  paste0("http://www.aejaffe.com/winterR_2017/",  
        "data/Charm_City_Circulator_Ridership.csv")  
)  
head(circ, 2)
```

A tibble: 2 × 15

	day	date	orangeBoardings	orangeAlightings	orangeAverage
	<chr>	<chr>	<int>	<int>	<dbl>
1	Monday	01/11/2010	877	1027	952
2	Tuesday	01/12/2010	777	815	796

... with 10 more variables: purpleBoardings <int>,
purpleAlightings <int>, purpleAverage <dbl>, greenBoardings <int>,
greenAlightings <int>, greenAverage <dbl>, bannerBoardings <int>,
bannerAlightings <int>, bannerAverage <dbl>, daily <dbl>

Creating a Date class from a character date

```
library(lubridate) # great for dates!  
library(dplyr) # mutate/summarise functions
```

Creating a Date class from a character date

```
sum(is.na(circ$date))
```

```
[1] 0
```

```
sum( circ$date == "")
```

```
[1] 0
```

```
circ = mutate(circ, date = mdy(date))  
sum( is.na(circ$date) ) # all converted correctly
```

```
[1] 0
```

```
head(circ$date, 3)
```

```
[1] "2010-01-11" "2010-01-12" "2010-01-13"
```

```
class(circ$date)
```


Reshaping data from wide (fat) to long (tall): base R

The `reshape` command exists. It is a **confusing** function. Don't use it.

tidyr package

tidyr allows you to “tidy” your data. We will be talking about:

- ▶ `gather` - make multiple columns into variables, (wide to long)
- ▶ `spread` - make a variable into multiple columns, (long to wide)
- ▶ `separate` - string into multiple columns
- ▶ `unite` - multiple columns into one string

Reshaping data from wide (fat) to long (tall): tidyr

`tidyr::gather` - puts column data into rows.

We want the column names into “var” variable in the output dataset and the value in “number” variable. We then describe which columns we want to “gather:”

```
library(tidyr)
long = gather(circ, key = "var", value = "number",
              starts_with("orange"), starts_with("purple"),
              starts_with("green"), starts_with("banner"))
head(long, 4)
```

A tibble: 4 × 5

	day	date	daily	var	number
	<chr>	<date>	<dbl>	<chr>	<dbl>
1	Monday	2010-01-11	952.0	orangeBoardings	877
2	Tuesday	2010-01-12	796.0	orangeBoardings	777
3	Wednesday	2010-01-13	1211.5	orangeBoardings	1203
4	Thursday	2010-01-14	1213.5	orangeBoardings	1194

Reshaping data from wide (fat) to long (tall): tidyr

```
table(long$var)
```

bannerAlightings	bannerAverage	bannerBoardings	greenA
1146	1146	1146	
greenAverage	greenBoardings	orangeAlightings	orang
1146	1146	1146	
orangeBoardings	purpleAlightings	purpleAverage	purple
1146	1146	1146	

Making a separator

We will use `str_replace` from `stringr` to put periods in the names:

```
library(stringr)
long = long %>% mutate(
  var = var %>% str_replace("Board", ".Board") %>%
    str_replace("Alight", ".Alight") %>%
    str_replace("Average", ".Average")
)
table(long$var)
```

banner.Alightings	banner.Average	banner.Boardings	green.Boardings
1146	1146	1146	1146
green.Average	green.Boardings	orange.Alightings	orange.Boardings
1146	1146	1146	1146
orange.Boardings	purple.Alightings	purple.Average	purple.Boardings
1146	1146	1146	1146

Reshaping data from wide (fat) to long (tall): tidyr

Now each var is boardings, averages, or alightings. We want to separate these so we can have these by line. Remember "." is special character:

```
long = separate(long, var, into = c("line", "type"),
                 sep = "[.]")
head(long, 2)
```

```
# A tibble: 2 × 6
```

	day <chr>	date <date>	daily <dbl>	line <chr>	type <chr>	number <dbl>
1	Monday	2010-01-11	952	orange	Boardings	877
2	Tuesday	2010-01-12	796	orange	Boardings	777

```
unique(long$line)
```

```
[1] "orange" "purple" "green"  "banner"
```

```
unique(long$type)
```

Re-uniting all the lines

If we had the opposite problem, we could use the `unite` function:

```
reunited = long %>%  
  unite(col = var, line, type, sep = ".")  
reunited %>% select(day, var) %>% head(3) %>% print
```

```
# A tibble: 3 × 2  
      day      var  
  <chr>    <chr>  
1 Monday orange.Boardings  
2 Tuesday orange.Boardings  
3 Wednesday orange.Boardings
```

We could also use `paste/paste0`.

Making column names a little more separated

Alternative: We could have replaced the column names first **then** reshaped:

```
cn = colnames(circ)
cn = cn %>%
  str_replace("Board", ".Board") %>%
  str_replace("Alight", ".Alight") %>%
  str_replace("Average", ".Average")
colnames(circ) = cn # then reshape using gather!
```


Reshaping data from long (tall) to wide (fat): tidyr

In tidyr, the spread function spreads rows into columns. Now we have a long data set, but we want to separate the Average, Alightings and Boardings into different columns:

```
# have to remove missing days
```

```
wide = filter(long, !is.na(date))  
wide = spread(wide, type, number)  
head(wide)
```

```
# A tibble: 6 × 7
```

	day	date	daily	line	Alightings	Average	Boardings
	<chr>	<date>	<dbl>	<chr>	<dbl>	<dbl>	<dbl>
1	Friday	2010-01-15	1644.0	banner	NA	NA	NA
2	Friday	2010-01-15	1644.0	green	NA	NA	NA
3	Friday	2010-01-15	1644.0	orange	1643	1644	1644
4	Friday	2010-01-15	1644.0	purple	NA	NA	NA
5	Friday	2010-01-22	1394.5	banner	NA	NA	NA
6	Friday	2010-01-22	1394.5	green	NA	NA	NA

Reshaping data from long (tall) to wide (fat): tidy

We can use `rowSums` to see if any values in the row is NA and keep if the row, which is a combination of date and line type has any non-missing data.

```
# wide = wide %>%  
#   select(Alightings, Average, Boardings) %>%  
#   mutate(good = rowSums(is.na(.)) > 0)  
not_namat = !is.na(select(wide, Alightings, Average, Boardings))  
head(not_namat, 2)
```

	Alightings	Average	Boardings
1	FALSE	FALSE	FALSE
2	FALSE	FALSE	FALSE

```
wide$good = rowSums(not_namat) > 0
```

Reshaping data from long (tall) to wide (fat): tidy

Now we can filter only the good rows and delete the good column.

```
wide = filter(wide, good) %>% select(-good)
head(wide)
```

```
# A tibble: 6 × 7
```

	day	date	daily	line	Alightings	Average	Boardings
	<chr>	<date>	<dbl>	<chr>	<dbl>	<dbl>	<dbl>
1	Friday	2010-01-15	1644.0	orange	1643	1644.0	1645
2	Friday	2010-01-22	1394.5	orange	1388	1394.5	1401
3	Friday	2010-01-29	1332.0	orange	1322	1332.0	1343
4	Friday	2010-02-05	1217.5	orange	1204	1217.5	1231
5	Friday	2010-02-12	671.0	orange	678	671.0	664
6	Friday	2010-02-19	1642.0	orange	1647	1642.0	1637

Finding the First (or Last) record

```
long = long %>% filter(!is.na(number) & number > 0)
first_and_last = long %>% arrange(date) %>% # arrange by date
  filter(type %in% "Boardings") %>% # keep boardings only
  group_by(line) %>% # group by line
  slice( c(1, n())) # select ("slice") first and last (n())
first_and_last %>% head(4)
```

Source: local data frame [4 x 6]

Groups: line [2]

	day	date	daily	line	type	number
	<chr>	<date>	<dbl>	<chr>	<chr>	<dbl>
1	Monday	2012-06-04	13342.5	banner	Boardings	520
2	Friday	2013-03-01	NA	banner	Boardings	817
3	Tuesday	2011-11-01	8873.0	green	Boardings	887
4	Friday	2013-03-01	NA	green	Boardings	2592

Data Merging/Append in Base R

- ▶ Merging - joining data sets together - usually on key variables, usually “id”
- ▶ `merge()` is the most common way to do this with data sets
- ▶ `rbind/cbind` - row/column bind, respectively
 - ▶ `rbind` is the equivalent of “appending” in Stata or “setting” in SAS
 - ▶ `cbind` allows you to add columns in addition to the previous ways
- ▶ `t()` is a function that will transpose the data

Merging

```
base <- data.frame(id = 1:10, Age= seq(55,60, length=10))  
base[1:2,]
```

	id	Age
1	1	55.00000
2	2	55.55556

```
visits <- data.frame(id = rep(1:8, 3), visit= rep(1:3, 8),  
                     Outcome = seq(10,50, length=24))  
visits[1:2,]
```

	id	visit	Outcome
1	1	1	10.00000
2	2	2	11.73913

Merging

```
merged.data <- merge(base, visits, by="id")  
merged.data[1:5,]
```

	id	Age	visit	Outcome
1	1	55.00000	1	10.00000
2	1	55.00000	3	23.91304
3	1	55.00000	2	37.82609
4	2	55.55556	2	11.73913
5	2	55.55556	1	25.65217

```
dim(merged.data)
```

```
[1] 24  4
```

Merging

```
all.data <- merge(base, visits, by="id", all=TRUE)
tail(all.data)
```

	id	Age	visit	Outcome
21	7	58.33333	2	48.26087
22	8	58.88889	2	22.17391
23	8	58.88889	1	36.08696
24	8	58.88889	3	50.00000
25	9	59.44444	NA	NA
26	10	60.00000	NA	NA

```
dim(all.data)
```

```
[1] 26  4
```


Joining in dplyr

- ▶ `?join` - see different types of joining for dplyr
- ▶ Let's look at <https://www.rstudio.com/wp-content/uploads/2015/02/data-wrangling-cheatsheet.pdf>

Left Join

```
lj = left_join(base, visits)
```

Joining, by = "id"

```
dim(lj)
```

```
[1] 26  4
```

```
tail(lj)
```

	id	Age	visit	Outcome
21	7	58.33333	2	48.26087
22	8	58.88889	2	22.17391
23	8	58.88889	1	36.08696
24	8	58.88889	3	50.00000
25	9	59.44444	NA	NA
26	10	60.00000	NA	NA

Right Join

```
rj = right_join(base, visits)
```

Joining, by = "id"

```
dim(rj)
```

```
[1] 24  4
```

```
tail(rj)
```

	id	Age	visit	Outcome
19	3	56.11111	1	41.30435
20	4	56.66667	2	43.04348
21	5	57.22222	3	44.78261
22	6	57.77778	1	46.52174
23	7	58.33333	2	48.26087
24	8	58.88889	3	50.00000

Full Join

```
fj = full_join(base, visits)
```

Joining, by = "id"

```
dim(fj)
```

```
[1] 26  4
```

```
tail(fj)
```


	id	Age	visit	Outcome
21	7	58.33333	2	48.26087
22	8	58.88889	2	22.17391
23	8	58.88889	1	36.08696
24	8	58.88889	3	50.00000
25	9	59.44444	NA	NA
26	10	60.00000	NA	NA

Bonus slides - explore after visualization!

Perform Operations By Multiple Groups: dplyr

We can then easily plot each day over time:


```
library(ggplot2)
ggplot(aes(x = date, y = Average, colour = line), data = w)
```



ggplot2 with piping

```
wide %>%
```

```
  ggplot(aes(x = date, y = Average, colour = line)) + geom_line()
```



Perform Operations By Multiple Groups: dplyr

Let's create the middle of the month (the 15th for example), and name it mon.

```
mon = wide %>%  
  mutate(month = month(date), year = year(date)) %>%  
  dplyr::group_by(line, month, year) %>%  
  dplyr::summarise(mean_avg = mean(Average))  
mon = mutate(mon, mid_month = dmy(paste0("15-", month, "-")),  
head(mon)
```

Source: local data frame [6 x 5]


Groups: line, month [6]

	line	month	year	mean_avg	mid_month
	<chr>	<dbl>	<dbl>	<dbl>	<date>
1	banner	1	2013	610.3226	2013-01-15
2	banner	2	2013	656.4643	2013-02-15
3	banner	3	2013	822.0000	2013-03-15
4	banner	6	2013	1000.1000	2013-06-15

Perform Operations By Multiple Groups: dplyr

We can then easily plot the mean of each month to see a smoother output:

```
ggplot(aes(x = mid_month,  
           y = mean_avg,  
           colour = line), data = mon) + geom_line()
```



Bonus! Points with a smoother!

```
ggplot(aes(x = date, y = Average, colour = line),  
       data = wide) + geom_smooth(se = FALSE) +  
       geom_point(size = .5)
```

`geom_smooth()` using method = 'gam'

